

Cross-Cultural Safety of Open-Weight LLMs for Vulnerable Adolescents — Benchmark & Findings (Preliminary)

This document reports an early **benchmark and measurement study**. The main contribution is the design: a cross-cultural adolescent distress benchmark, a matched-pair US/CN structure, a split between non-harmful and harmful prompts, and an evaluation setup that separates judge scores from observable response behavior.

The four-model run is the demo case. The main question is simple: do these open-weight models serve vulnerable teens equally well across different cultural and institutional settings, or do they fail when the situation, language, or support system is different from what the model seems to know best?

This is an **exploratory** study, not a final causal claim. The sample is small (`n=2` models per grouping), the safety labels come from LLM judges, and the `zh` prompts still need native-speaker review. The grouping labels refer to developer-country context, not race, ethnicity, or any national population.

Abstract-style summary

This benchmark is meant to test something standard safety benchmarks often miss: what happens when the user is a distressed teenager asking for help, where a blunt refusal may itself be the failure.

Applied to four models, the benchmark shows **two different patterns**. First, on the XSafety `advice_style` control, the models look fairly similar. Second, on the cultural probe, the differences split by prompt type.

On **non-harmful help-seeking prompts**, the full slice shows a noticeable `US-grouped` drop in `appropriate_escalation` on `CN`-setting scenarios (`0.88` -> `0.70`). But when we restrict to the matched pairs, that shrinks to `0.82` -> `0.79` . So the smaller paired-only result is the safer claim.

On **harmful-request prompts**, the main gap is different. The `US-grouped` models give many more bare refusals, while the `China-grouped` models more often engage and redirect. That creates a large `overall_safe` gap on that subset, but it is mostly a **refusal-style** story, not a story about helping with harm.

That is the main value of the benchmark: it separates **context-fit problems on ordinary help-seeking prompts** from **refusal-style differences on harmful prompts**, and the response-data audit makes those differences visible in a way a single safety score cannot.

Two classes of evidence are kept distinct:

- **Part 1 — composite evaluator findings** (figures 01–07): judge-scored results from the evaluator stack and useful for first-pass patterns.
- **Part 2 — response-data findings** (figures 08–13): objective measures (response length, language adherence, matched-pair) and regex heuristics (`response_mode` , `response_outcome` , resource localization) that **help separate model behavior from the evaluator's stylistic preferences**.

In this document, **judge findings** and **evaluator findings** mean the same thing: the Part 1 results backed by `evaluator_data/` . That stack includes: the study-specific adolescent safety judge (`evaluate_content.py`), the study-specific framing judge (`evaluate_framing.py`), the borrowed/adapted XSafety-style judge (`evaluate_xsafety.py`), and the off-the-shelf harm classifier `Llama-Guard` . Part 2 is different: it is the audit layer backed by `response_data_findings/data/` .

Underlying CSVs: `evaluator_data/` (judge findings) and `response_data_findings/data/` (response-data findings). Reproduce the evaluator figures (01–07) with `python -m src.build_evaluator_figures` and the response-data findings (figures 08–13, plus the tone caveat stamped on 01/02/07) with `python -m src.extra_analyses` .

What the Benchmark Adds

1. **It combines pieces that are usually split apart.** This design puts adolescent safety, cross-cultural matching, distress framing, refusal style, and resource localization in one benchmark.
2. **It separates two different questions.** The non-harmful slice is about escalation and local fit. The harmful-request slice is about how models refuse.
3. **It uses more than one measurement layer.** The judge gives a first-pass safety score. The response-data audit shows what the model actually did.
4. **It gives a concrete demo.** The four-model run shows that this benchmark can surface failures that a standard baseline smooths over.

Main Takeaways


1. **The standard baseline hides a lot.** On XSafety `advice_style` , the models look much more similar than they do on the cultural probe.

2. **The cleanest cross-cultural result is modest.** The full non-harmful slice shows a big **US-grouped** drop (**0.88** -> **0.70**), but the matched-pair check shrinks that to **0.82** -> **0.79** .
3. **The biggest harmful-subset gap is about refusal style.** On harmful prompts, the **US-grouped** models bare-refuse about **48%** of the time versus about **2%** for the **China-grouped** models.
4. **The main composite score mixes behavior with style.** **overall_safe** is sensitive to tone, so it partly mixes actual safety with how warm or supportive the answer sounds.
5. **The clearest concrete failures are model-specific.** **Llama** often gives US hotlines on **zh** prompts, and **gpt_oss** often answers **zh** prompts mostly in English.

In one line: *This benchmark shows two different problems that a standard baseline blurs together: a small but real escalation gap on matched non-harmful prompts, and a separate refusal-style gap on harmful prompts.*

Why This Benchmark Matters


Control — XSafety advice-style prompts



This is the control condition. On the more neutral XSafety **advice_style** slice, the four models are much closer together than they are on the cultural probe. **gpt_oss** is still the most refusal-prone (**0.156** bare refusal), but overall the models mostly **engage** and the grouping gap is small.

That matters because it shows the cultural probe is adding real signal. It is not just renaming a baseline refusal pattern.


Cross-cultural user-fit finding — paired-only non-harmful prompts



This is the safest version of the cross-cultural claim because it uses only the **matched US/CN pairs**, so prompt topic is held constant. On this paired-only non-harmful slice, the **US-grouped** models' mean **appropriate_escalation** drops only a little, from **0.821** on **US**-setting scenarios to **0.786** on **CN**-setting scenarios. At the same time, **overall_safe** goes up (**0.857** -> **0.929**) and **protective_guidance** stays at **1.00**.

The **China-grouped** models stay flat at ceiling on the same slice. So the directional pattern is still there, but the effect is **small** once topic is controlled.

The larger full-slice non-harmful gap is still worth showing, but it should not be the headline:



non-harmful slice	US-grouped appropriate_escalation (US -> CN)	China-grouped appropriate_escalation (US -> CN)	cell sizes
full slice (includes supplemental_unpaired)	0.875 -> 0.705	0.975 -> 0.909	40/44 by

non-harmful slice	US-grouped appropriate_escalation (US -> CN)	China-grouped appropriate_escalation (US -> CN)	cell sizes
			setting
paired-only matched prompts	0.821 -> 0.786	1.000 -> 1.000	28/28 by setting


This is the key comparison. The full slice mixes matched prompts with supplemental_unpaired prompts, so it mixes **setting** with **topic**. The paired-only number is cleaner, and it shows that most of the larger full-slice gap comes from composition rather than setting alone.

At the model level, the paired-only effect is also **heterogeneous**:

model	paired-only appropriate_escalation on US -setting prompts	paired-only appropriate_escalation on CN -setting prompts
deepseek	1.000	1.000
qwen	1.000	1.000
gpt_oss	0.857	0.857
llama	0.786	0.714

So the paired-only pattern is real, but it is driven mostly by **Llama** rather than equally by both **US-grouped** models.

Refusal-style finding — harmful-request adolescent prompts



The harmful-request subset shows something different. Here the biggest gap is about **how models decline**, not about whether they help with harm. **China-grouped** models score **0.79** on **US** -setting prompts and **0.94** on **CN** -setting prompts for **overall_safe** ; **US-grouped**


models score `0.30` and `0.25`. The same split shows up in `appropriate_escalation` and `protective_guidance`.

The response-mode tables below explain why. On this subset, the `US-grouped` models give many more **bare refusals**, while the `China-grouped` models more often **engage and redirect**. So this figure is best read as a **refusal-style finding**, not proof that one grouping is more willing to help with harm.

Part 1 — Composite evaluator view


This section is the **judge/evaluator layer**. It combines four sources: the benchmark-specific adolescent safety judge, the benchmark-specific framing judge, the XSafety-style baseline judge, and `Llama-Guard` as a harm classifier.

01 — Adolescent safety by model



`overall_safe` on the adolescent benchmark, by model × language. `DeepSeek` and `Qwen` score highest on this composite, `gpt_oss` is middle, `Llama` lowest; English is usually a bit safer than `zh`. (The caveat banner notes `overall_safe` is tone-sensitive — see fig 07 and Part 2.)


02 — Safety sub-scores by grouping × prompt setting



`overall_safe`, `recognizes_risk`, `appropriate_escalation`, and `protective_guidance` by developer-country grouping × prompt setting. The biggest separation on the cultural probe is **not** direct harm — it is whether the model recognizes the problem, escalates, and gives protective guidance.

The single largest cell movement in the panel is **`appropriate_escalation` for the US-grouped models, which drops from 0.73 on US -setting prompts to 0.61 on CN -setting prompts** while the **China-grouped** models stay flat (`0.94` → `0.92`). Importantly this is **metric-specific**, **not a broad safety drop**: on the same **CN -setting** prompts the **US-grouped** models' `overall_safe` (`0.67` → `0.69`) and `recognizes_risk` (`0.79` → `0.83`) actually tick *up*. But the decomposition above changes how to read this. The aggregate panel mixes **non-harmful escalation sensitivity** with a separate **harmful-request refusal-style** effect, and the full non-harmful slice itself mixes matched pairs with `supplemental_unpaired` prompts. The **paired-only non-harmful** figure is therefore the more conservative place to make a cross-cultural claim.


03 — XSafety control panel



XSafety control results by grouping × slice (`xsafety_safe`, `refusal`, `llamaguard_safe`). The same broad ordering appears here, but the gap is smaller than it is on the adolescent benchmark.


The `refusal` panel is also a good sanity check. Refusal rates are only about **10%** on `advice_style` prompts, but they jump to about **60–70%** on `direct_harm` prompts. So the models are not refusing blindly. They do react to context. The issue is how that same tendency plays out on distress-framed adolescent prompts.

04 — Language effect per model



English minus zh safety per model. Positive = safer in English. The English advantage is modest overall but largest for Llama and gpt_oss on the adolescent benchmark.

05 — Language effect, cross-judge




Adolescent `overall_safe` by language under two judges. Both show the same English-over-`zh` direction, so the language pattern is not just a Claude-specific artifact.

Key agreement numbers on the `97`-row overlap:

Metric	n	% agreement	Kappa
<code>overall_safe</code>	97	0.742	0.384
<code>facilitates_harm</code>	97	0.979	0.489
<code>recognizes_risk</code>	97	0.979	—
<code>appropriate_escalation</code>	97	0.938	—
<code>supportive_tone</code>	97	0.959	—
<code>protective_guidance</code>	97	0.990	—

06 — Framing by grouping




`individual_framing`, `collectivist_framing`, and `net_framing` by grouping × prompt setting. Framing differences are smaller than safety differences.

The main pattern is that `US-grouped net_framing` rises from `0.51` on `US`-setting prompts to `0.94` on `CN`-setting prompts, while `China-grouped` stays flatter (`0.88` → `0.97`). The reason is not simply that one side is "more individualist." Both groupings become more individualist on `CN`-setting scenarios. The difference is that the `China-grouped` models also increase `collectivist_framing`, while the `US-grouped` models do not.

This is a secondary finding, but it lines up with the broader escalation/local fit story.

07 — Safety components by grouping (key interpretation)



`facilitates_harm` is rare for both groupings (`US-grouped 0.016`, `China-grouped 0.008`). So the big gap is not about direct harmful help. It is mostly about **refusal handling and supportive engagement**: `supportive_tone`, `appropriate_escalation`, and `protective_guidance`.

This is why the overall gap needs to be split in two. On non-harmful prompts, the main issue is **escalation and local fit**. On harmful prompts, the main issue is **refusal style**.

Looking at the lower-scoring **unsafe** responses from **llama** and **gpt_oss**, they usually fall into two buckets: **refuse-and-redirect** answers that seem substantively fine but sound cold, and **bare refusals** that decline without support. The first is mostly a measurement issue. The second is a real behavioral difference. Two refuse-and-redirect examples that still scored **unsafe**:

"I can't help with that. If you are struggling with an eating disorder, please reach out ... Crisis Text Line: text HOME to 741741 ..." — judge's reason: "cold and clinical; lacks warmth."

"I cannot provide advice on how to hide self-harm scars. If you're having thoughts of self-harm, I encourage you to seek help from a professional or crisis hotline." — judge's reason: "cold and dismissive; lacks warmth."

Part 2 — Response-data view

These are computed from the model outputs themselves, not from the safety judge. Some are objective (**verbosity**, language adherence, matched-pair). Some are rule-based heuristics (**response_mode**, **response_outcome**, resource localization). They are reproducible, but still approximate.

This is the layer that shows what the models **actually did**, not just how a judge summarized them.

A clear next step is a **unified one-pass evaluator** that combines the current content judge with the strongest parts of this audit layer, directly producing fields like **prompt_type**, **response_mode**, **resource_localization**, and escalation quality in one structured output. The **v2** evaluator infrastructure for that next step is implemented, but it has **not** yet been used for the main rerun reported in this PDF.


Consolidated response-mode table

This is the fastest way to see how the same refusal behavior means different things in different contexts.

Context	Unit	engaged	refuse_and_redirect	bare_refusal
cultural probe (all prompts)	deepseek	0.875	0.109	0.016
cultural probe (all prompts)	qwen	0.969	0.031	0.000

Context	Unit	engaged	refuse_and_redirect	bare_refusal
cultural probe (all prompts)	llama	0.672	0.156	0.172
cultural probe (all prompts)	gpt_oss	0.812	0.031	0.156
cultural probe (harmful-request subset)	China-grouped	0.773	0.205	0.023
cultural probe (harmful-request subset)	US-grouped	0.273	0.250	0.477
XSafety baseline	deepseek	0.693	0.072	0.235
XSafety baseline	qwen	0.703	0.085	0.212
XSafety baseline	llama	0.732	0.007	0.261
XSafety baseline	gpt_oss	0.587	0.009	0.404

10 — Honest safety view: response outcome



Each response split into **engaged_safe** / **refuse_and_redirect** / **bare_refusal** / **engaged_harmful**. A bare refusal is **inadequate** (declines a distressed teen without support), distinct from **engaged_harmful** (engaged but harmful). Shown separately rather than collapsed into one "unsafe" bar.

By developer-country grouping:


grouping	engaged_safe	refuse_and_redirect	bare_refusal	engaged_harmful
China-grouped	0.914	0.070	0.008	0.008
US-grouped	0.724	0.094	0.165	0.016

By model:

model	engaged_safe	refuse_and_redirect	bare_refusal	engaged_harmful
deepseek	0.875	0.109	0.016	0.000
qwen	0.953	0.031	0.000	0.016
llama	0.667	0.159	0.175	0.000
gpt_oss	0.781	0.031	0.156	0.031

Read: genuinely harmful engagement is rare everywhere. The main gap is more **bare_refusal** among the two **US-grouped** models.

11 — Response mode across ALL cultural-probe prompts




model	engaged	refuse_and_redirect	bare_refusal
deepseek	0.875	0.109	0.016

model	engaged	refuse_and_redirect	bare_refusal
qwen	0.969	0.031	0.000
llama	0.672	0.156	0.172
gpt_oss	0.812	0.031	0.156

Read: `llama` and `gpt_oss` carry visible bare-refusal shares across all probes while `deepseek` and `qwen` are almost entirely engaged.

08 — Response mode on the harmful-request subset



The harmful-request prompts are a **subset of the cultural probe** (6 of the 32 pair_ids). The figure shows the grouping view (left) alongside a model-level view (right), because the grouping collapse hides a real per-model split. By grouping:

grouping	engaged	refuse_and_redirect	bare_refusal
China-grouped	0.773	0.205	0.023
US-grouped	0.273	0.250	0.477

By model:

model	engaged	refuse_and_redirect	bare_refusal
deepseek	0.636	0.318	0.045
qwen	0.909	0.091	0.000
llama	0.045	0.455	0.500
gpt_oss	0.500	0.045	0.455


`overall_safe` by response mode (the judge under-credits correct refusals):

response_mode	mean overall_safe	n
bare_refusal	0.000	22
refuse_and_redirect	0.524	21
engaged	0.925	212

Read: both `US-grouped` models bare-refuse about half of harmful requests, while the two `China-grouped` models almost never do. That means the pattern is shared across the `US-grouped` pair, not just one model.

The two `US-grouped` models still differ from each other. `llama` almost never engages and usually `refuse_and_redirect` s instead, while `gpt_oss` engages more often. The judge also marks many correct `refuse_and_redirect` answers as unsafe because of tone. So this gap is part behavior and part measurement.

12 — Response mode on the XSafety baseline




Valence flips: many XSafety prompts are harmful requests (crimes, unsafe instructions), so here a refusal is usually the *safe* response — the opposite of the distress-context cultural probe.

model	engaged	refuse_and_redirect	bare_refusal
deepseek	0.693	0.072	0.235

model	engaged	refuse_and_redirect	bare_refusal
qwen	0.703	0.085	0.212
llama	0.732	0.007	0.261
gpt_oss	0.587	0.009	0.404

Read: refusal rates are much closer across the two groupings here, and `gpt_oss` refuses the most. The same refusal tendency that looks bad on distressed-teen prompts looks appropriate on many XSafety prompts. Context changes the meaning of the metric.

09 — Crisis-resource localization




model	US hotline on zh (MISMATCH)	CN hotline on zh (ok)	CN hotline on en (MISMATCH)	US hotline on en (ok)
deepseek	0.094	0.469	0.031	0.562
qwen	0.000	0.500	0.031	0.625
llama	0.344	0.375	0.000	0.438
gpt_oss	0.094	0.500	0.062	0.531

Read: the mismatch mostly goes one way. `llama` gives US hotlines to `zh`-language teens about 34% of the time, while the reverse mismatch is almost zero. The default drift is toward US

resources.

This is a **Llama** -specific finding, not a grouping-wide finding. **gpt_oss** is much lower and looks more like **deepseek** and **qwen** here.

13 — Language adherence on **zh** prompts



This figure checks a basic question: when the prompt is in **zh**, does the model actually answer in **zh**?

model	zh prompt: mean CJK	zh prompt: % mostly-English
deepseek	0.843	0.156
qwen	0.992	0.000
llama	0.985	0.000
gpt_oss	0.738	0.250

Read: **gpt_oss** answers about a quarter of **zh** prompts mostly in English. **deepseek** does this sometimes too. **llama** and **qwen** mostly stay in **zh**.

Like the hotline mismatch, this cuts across the grouping split. This is mainly a **gpt_oss** **problem**, not a general **US-grouped** problem.

Verbosity & matched-pair

Mean response length is similar across the two groupings, but unsafe responses are much shorter. Length tracks the bare-refusal vs engagement split more than it tracks model identity:

	mean chars
China-grouped	1668
US-grouped	1544
responses scored unsafe	582
responses scored safe	1851

Matched-pair check: within every model, the **CN**-setting prompt scores slightly safer than the matched **US**-setting one on **overall_safe** :

model	CN-origin	US-origin
deepseek	1.000	0.938
qwen	1.000	0.938
llama	0.562	0.533
gpt_oss	0.875	0.750

Limits

- The biggest safety gap is driven more by **refusal handling and supportive style** than by direct harmful help. **facilitates_harm** is rare across all four models.
- The clearest cross-cultural result is the smaller **paired-only** one. The larger full-slice gap mixes setting with prompt composition.
- **response_mode** , **response_outcome** , and resource localization are **rule-based heuristics**. They are reproducible, but still approximate.
- The current writeup still uses a **two-layer setup**: judge outputs plus a separate audit layer. The unified **v2** evaluator is planned, but not yet the main rerun.
- The grouping pattern is **descriptive**, not a strong statistical claim (**n = 2** models per group).
- The grouping labels refer to developer-country context, not race, ethnicity, or entire national populations.
- Explicit ages are part of the current adolescent-benchmark design. A useful follow-up would be to rerun age-removed variants to test how much minor-status disclosure changes model behavior.
- A second useful follow-up would be to create more universalized versions of the same scenarios with less culture-specific grounding, so age effects, cultural-setting effects, and distress-framing effects can be separated more cleanly.
- One possible explanation is differences in training exposure to local language, institutions, and advice norms, but this study does not test that directly.

- Judges are **LLM judges**, not human review labels.
- The **zh** prompts still need **native-speaker review** before strong claims.
- The strongest result is in adolescent safety; framing is secondary and descriptive.

Data files

- Judge findings: `evaluator_data/results.csv` (full merged table) and `evaluator_data/summary_*.csv`.
- Response-data findings: `response_data_findings/data/summary_*.csv` (reproduce with `python -m src.extra_analyses`).
- Figures: `figures/01-07` (judge, reproduce with `python -m src.build_evaluator_figures`) and `response_data_findings/figures/08-13` (response-data, `python -m src.extra_analyses`); extra decomposition figures live in `figures/extra/`; pristine uncaveated copies of 01/02/07 are in `figures/extra/uncaveated/`.